

Classification

March 15, 2022

1 Supervised machine learning: Classification

Dataset Introduction The dataset used here is Titanic - Machine Learning from Disaster. This dataset is used for an example for classification problems. The datasets consists of the following columns

Features

survival - This feature tells us whether the passenger survived or not. This feature is gonna be used for predicting 0 = No; 1 = Yes

PassengerId - ID of the observation

pclass Passenger Class (1 = 1st; 2 = 2nd; 3 = 3rd)

name - Name of the passenger

sex- Sex of the passenger

Age- Age of the passenger

sibsp - Number of siblings or spouse traveling with the passenger

parch - Number of parents, children traveling with the passenger

ticket- Ticket number of the passenger

fare - Fare amount paid by the passenger

cabin- Cabin allocated for the passenger

embarked - Place of embarkment of the passenger (C = Cherbourg; Q = Queenstown; S = Southampton)

Problem Statement - Predict survivability of a passenger based on the features The models will be focus on prediction rather than interpretability

```
[5]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import LogisticRegression
```

```

from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedShuffleSplit
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import MinMaxScaler
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import precision_recall_fscore_support as score

from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix, accuracy_score, roc_auc_score
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import average_precision_score
from sklearn.metrics import roc_curve
from sklearn.metrics import auc
from sklearn.model_selection import cross_val_score
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report

%matplotlib inline

```

```

[1]: from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```

[2]: import seaborn as sns

```

```

[6]: data = pd.read_csv('/content/drive/MyDrive/ibm/Projects/EDA/train(EDA).csv')

```

Initial exploration of data

```

[7]: print(data.shape)
print(data.columns.tolist())
print(data.dtypes)

```

```

(891, 12)
['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch',
'Ticket', 'Fare', 'Cabin', 'Embarked']
PassengerId      int64
Survived          int64
Pclass            int64

```

```
Name          object
Sex           object
Age          float64
SibSp        int64
Parch        int64
Ticket       object
Fare         float64
Cabin        object
Embarked     object
dtype: object
```

```
[8]: data.head()
```

```
[8]:   PassengerId  Survived  Pclass  \
0             1         0        3
1             2         1        1
2             3         1        3
3             4         1        1
4             5         0        3
```

```
                                Name    Sex  Age  SibSp  \
0                                Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2                                Heikkinen, Miss. Laina  female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)    female  35.0      1
4                                Allen, Mr. William Henry    male  35.0      0
```

```
   Parch    Ticket   Fare Cabin Embarked
0      0  A/5 21171   7.2500   NaN        S
1      0    PC 17599  71.2833   C85        C
2      0 STON/O2. 3101282   7.9250   NaN        S
3      0    113803  53.1000  C123        S
4      0    373450   8.0500   NaN        S
```

We can see that the problem to solve here is to predict the survived column

```
[9]: data.describe()
```

```
[9]:   PassengerId  Survived  Pclass    Age  SibSp  \
count  891.000000  891.000000  891.000000  714.000000  891.000000
mean    446.000000    0.383838    2.308642   29.699118    0.523008
std     257.353842    0.486592    0.836071   14.526497    1.102743
min       1.000000    0.000000    1.000000    0.420000    0.000000
25%     223.500000    0.000000    2.000000   20.125000    0.000000
50%     446.000000    0.000000    3.000000   28.000000    0.000000
75%     668.500000    1.000000    3.000000   38.000000    1.000000
max     891.000000    1.000000    3.000000   80.000000    8.000000
```

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

2 Feature deletion

Finding columns which has unique values for each observation

```
[10]: columnname = data.columns.to_list()
observation_length = data.shape[0]
dummies = []
for x in columnname:
    if len(data[x].unique()) == observation_length:
        dummies.append(x)
print(dummies)
```

```
['PassengerId', 'Name']
```

As they do not provide any useful info. We are gonna delete them

```
[11]: data1 = data.copy()
data1.drop(dummies, axis = 1, inplace = True)
data1.shape
```

```
[11]: (891, 10)
```

```
[12]: data1.head()
```

```
[12]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	\
0	0	3	male	22.0	1	0	A/5 21171	7.2500	
1	1	1	female	38.0	1	0	PC 17599	71.2833	
2	1	3	female	26.0	0	0	STON/O2. 3101282	7.9250	
3	1	1	female	35.0	1	0	113803	53.1000	
4	0	3	male	35.0	0	0	373450	8.0500	

	Cabin	Embarked
0	NaN	S
1	C85	C
2	NaN	S
3	C123	S

4 NaN S

the ticket columns looks a little suspicious

```
[13]: data1['Ticket'].describe()
```

```
[13]: count          891
      unique         681
      top           347082
      freq           7
      Name: Ticket, dtype: object
```

it has a huge number of unique values lets check them to see any order or categories

```
[14]: data['Ticket'].unique()
```

```
[14]: array(['A/5 21171', 'PC 17599', 'STON/O2. 3101282', '113803', '373450',
        '330877', '17463', '349909', '347742', '237736', 'PP 9549',
        '113783', 'A/5. 2151', '347082', '350406', '248706', '382652',
        '244373', '345763', '2649', '239865', '248698', '330923', '113788',
        '347077', '2631', '19950', '330959', '349216', 'PC 17601',
        'PC 17569', '335677', 'C.A. 24579', 'PC 17604', '113789', '2677',
        'A./5. 2152', '345764', '2651', '7546', '11668', '349253',
        'SC/Paris 2123', '330958', 'S.C./A.4. 23567', '370371', '14311',
        '2662', '349237', '3101295', 'A/4. 39886', 'PC 17572', '2926',
        '113509', '19947', 'C.A. 31026', '2697', 'C.A. 34651', 'CA 2144',
        '2669', '113572', '36973', '347088', 'PC 17605', '2661',
        'C.A. 29395', 'S.P. 3464', '3101281', '315151', 'C.A. 33111',
        'S.O.C. 14879', '2680', '1601', '348123', '349208', '374746',
        '248738', '364516', '345767', '345779', '330932', '113059',
        'SO/C 14885', '3101278', 'W./C. 6608', 'SOTON/OQ 392086', '343275',
        '343276', '347466', 'W.E.P. 5734', 'C.A. 2315', '364500', '374910',
        'PC 17754', 'PC 17759', '231919', '244367', '349245', '349215',
        '35281', '7540', '3101276', '349207', '343120', '312991', '349249',
        '371110', '110465', '2665', '324669', '4136', '2627',
        'STON/O 2. 3101294', '370369', 'PC 17558', 'A4. 54510', '27267',
        '370372', 'C 17369', '2668', '347061', '349241',
        'SOTON/O.Q. 3101307', 'A/5. 3337', '228414', 'C.A. 29178',
        'SC/PARIS 2133', '11752', '7534', 'PC 17593', '2678', '347081',
        'STON/O2. 3101279', '365222', '231945', 'C.A. 33112', '350043',
        '230080', '244310', 'S.O.P. 1166', '113776', 'A.5. 11206',
        'A/5. 851', 'Fa 265302', 'PC 17597', '35851', 'SOTON/OQ 392090',
        '315037', 'CA. 2343', '371362', 'C.A. 33595', '347068', '315093',
        '363291', '113505', 'PC 17318', '111240', 'STON/O 2. 3101280',
        '17764', '350404', '4133', 'PC 17595', '250653', 'LINE',
        'SC/PARIS 2131', '230136', '315153', '113767', '370365', '111428',
        '364849', '349247', '234604', '28424', '350046', 'PC 17610',
        '368703', '4579', '370370', '248747', '345770', '3101264', '2628',
```

'A/5 3540', '347054', '2699', '367231', '112277',
'SOTON/O.Q. 3101311', 'F.C.C. 13528', 'A/5 21174', '250646',
'367229', '35273', 'STON/O2. 3101283', '243847', '11813',
'W/C 14208', 'SOTON/OQ 392089', '220367', '21440', '349234',
'19943', 'PP 4348', 'SW/PP 751', 'A/5 21173', '236171', '347067',
'237442', 'C.A. 29566', 'W./C. 6609', '26707', 'C.A. 31921',
'28665', 'SCD/W 1585', '367230', 'W./C. 14263',
'STON/O 2. 3101275', '2694', '19928', '347071', '250649', '11751',
'244252', '362316', '113514', 'A/5. 3336', '370129', '2650',
'PC 17585', '110152', 'PC 17755', '230433', '384461', '110413',
'112059', '382649', 'C.A. 17248', '347083', 'PC 17582', 'PC 17760',
'113798', '250644', 'PC 17596', '370375', '13502', '347073',
'239853', 'C.A. 2673', '336439', '347464', '345778', 'A/5. 10482',
'113056', '349239', '345774', '349206', '237798', '370373',
'19877', '11967', 'SC/Paris 2163', '349236', '349233', 'PC 17612',
'2693', '113781', '19988', '9234', '367226', '226593', 'A/5 2466',
'17421', 'PC 17758', 'P/PP 3381', 'PC 17485', '11767', 'PC 17608',
'250651', '349243', 'F.C.C. 13529', '347470', '29011', '36928',
'16966', 'A/5 21172', '349219', '234818', '345364', '28551',
'111361', '113043', 'PC 17611', '349225', '7598', '113784',
'248740', '244361', '229236', '248733', '31418', '386525',
'C.A. 37671', '315088', '7267', '113510', '2695', '2647', '345783',
'237671', '330931', '330980', 'SC/PARIS 2167', '2691',
'SOTON/O.Q. 3101310', 'C 7076', '110813', '2626', '14313',
'PC 17477', '11765', '3101267', '323951', 'C 7077', '113503',
'2648', '347069', 'PC 17757', '2653', 'STON/O 2. 3101293',
'349227', '27849', '367655', 'SC 1748', '113760', '350034',
'3101277', '350052', '350407', '28403', '244278', '240929',
'STON/O 2. 3101289', '341826', '4137', '315096', '28664', '347064',
'29106', '312992', '349222', '394140', 'STON/O 2. 3101269',
'343095', '28220', '250652', '28228', '345773', '349254',
'A/5. 13032', '315082', '347080', 'A/4. 34244', '2003', '250655',
'364851', 'SOTON/O.Q. 392078', '110564', '376564', 'SC/AH 3085',
'STON/O 2. 3101274', '13507', 'C.A. 18723', '345769', '347076',
'230434', '65306', '33638', '113794', '2666', '113786', '65303',
'113051', '17453', 'A/5 2817', '349240', '13509', '17464',
'F.C.C. 13531', '371060', '19952', '364506', '111320', '234360',
'A/S 2816', 'SOTON/O.Q. 3101306', '113792', '36209', '323592',
'315089', 'SC/AH Basle 541', '7553', '31027', '3460', '350060',
'3101298', '239854', 'A/5 3594', '4134', '11771', 'A.5. 18509',
'65304', 'SOTON/OQ 3101317', '113787', 'PC 17609', 'A/4 45380',
'36947', 'C.A. 6212', '350035', '315086', '364846', '330909',
'4135', '26360', '111427', 'C 4001', '382651', 'SOTON/OQ 3101316',
'PC 17473', 'PC 17603', '349209', '36967', 'C.A. 34260', '226875',
'349242', '12749', '349252', '2624', '2700', '367232',
'W./C. 14258', 'PC 17483', '3101296', '29104', '2641', '2690',
'315084', '113050', 'PC 17761', '364498', '13568', 'WE/P 5735',

```
'2908', '693', 'SC/PARIS 2146', '244358', '330979', '2620',
'347085', '113807', '11755', '345572', '372622', '349251',
'218629', 'SOTON/OQ 392082', 'SOTON/O.Q. 392087', 'A/4 48871',
'349205', '2686', '350417', 'S.W./PP 752', '11769', 'PC 17474',
'14312', 'A/4. 20589', '358585', '243880', '2689',
'STON/O 2. 3101286', '237789', '13049', '3411', '237565', '13567',
'14973', 'A./5. 3235', 'STON/O 2. 3101273', 'A/5 3902', '364848',
'SC/AH 29037', '248727', '2664', '349214', '113796', '364511',
'111426', '349910', '349246', '113804', 'SOTON/O.Q. 3101305',
'370377', '364512', '220845', '31028', '2659', '11753', '350029',
'54636', '36963', '219533', '349224', '334912', '27042', '347743',
'13214', '112052', '237668', 'STON/O 2. 3101292', '350050',
'349231', '13213', 'S.O./P.P. 751', 'CA. 2314', '349221', '8475',
'330919', '365226', '349223', '29751', '2623', '5727', '349210',
'STON/O 2. 3101285', '234686', '312993', 'A/5 3536', '19996',
'29750', 'F.C. 12750', 'C.A. 24580', '244270', '239856', '349912',
'342826', '4138', '330935', '6563', '349228', '350036', '24160',
'17474', '349256', '2672', '113800', '248731', '363592', '35852',
'348121', 'PC 17475', '36864', '350025', '223596', 'PC 17476',
'PC 17482', '113028', '7545', '250647', '348124', '34218', '36568',
'347062', '350048', '12233', '250643', '113806', '315094', '36866',
'236853', 'STON/O2. 3101271', '239855', '28425', '233639',
'349201', '349218', '16988', '376566', 'STON/O 2. 3101288',
'250648', '113773', '335097', '29103', '392096', '345780',
'349204', '350042', '29108', '363294', 'SOTON/O2 3101272', '2663',
'347074', '112379', '364850', '8471', '345781', '350047',
'S.O./P.P. 3', '2674', '29105', '347078', '383121', '36865',
'2687', '113501', 'W./C. 6607', 'SOTON/O.Q. 3101312', '374887',
'3101265', '12460', 'PC 17600', '349203', '28213', '17465',
'349244', '2685', '2625', '347089', '347063', '112050', '347087',
'248723', '3474', '28206', '364499', '112058', 'STON/O2. 3101290',
'S.C./PARIS 2079', 'C 7075', '315098', '19972', '368323', '367228',
'2671', '347468', '2223', 'PC 17756', '315097', '392092', '11774',
'SOTON/O2 3101287', '2683', '315090', 'C.A. 5547', '349213',
'347060', 'PC 17592', '392091', '113055', '2629', '350026',
'28134', '17466', '233866', '236852', 'SC/PARIS 2149', 'PC 17590',
'345777', '349248', '695', '345765', '2667', '349212', '349217',
'349257', '7552', 'C.A./SOTON 34068', 'SOTON/OQ 392076', '211536',
'112053', '111369', '370376'], dtype=object)
```

Ticket columns doesnt seem to have any ordering or category, As this also doesnt give any useful info. We shall delete it also

```
[15]: data2 = data1.copy()
data2.drop(['Ticket'], axis = 1, inplace = True)
data2.shape
```

```
[15]: (891, 9)
```

3 Exploratory data analysis and cleaning

3.1 Numerical columns Analysis

```
[16]: numerical_cols = data2.select_dtypes('number').columns.tolist()
      print(numerical_cols)
```

```
['Survived', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare']
```

```
[17]: data2.describe()
```

```
[17]:
```

	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

3.1.1 Survived Column

```
[18]: data2['Survived'].isnull().sum()
```

```
[18]: 0
```

```
[19]: data2['Survived'].unique()
```


```
[19]: array([0, 1])
```

```
[20]: ct = pd.crosstab(index = data2['Survived'], columns = 'count')
      ct = ct/data2.shape[0]
      display(ct)
      ct.plot(kind = 'bar')
```

```
col_0    count
Survived
0         0.616162
1         0.383838
```



```
[20]: <matplotlib.axes._subplots.AxesSubplot at 0x7f109681d190>
```



We can see that around 61 % people died and around 38% people survived

3.1.2 Pclass column

```
[21]: data2['Pclass'].isnull().sum()
```

```
[21]: 0
```

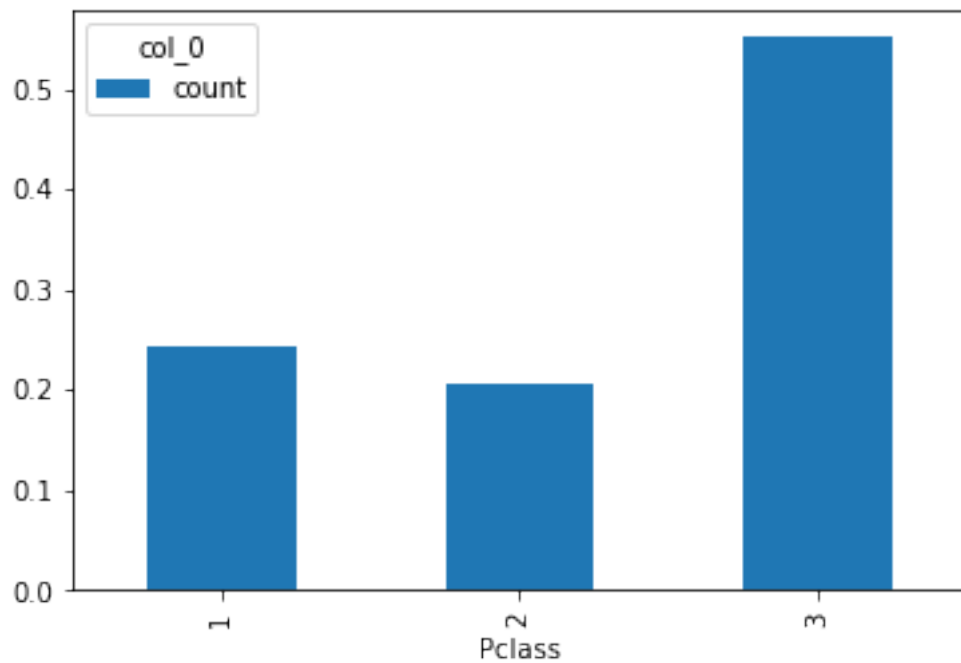
```
[22]: data2['Pclass'].unique()
```

```
[22]: array([3, 1, 2])
```

```
[23]: ct = pd.crosstab(index = data2['Pclass'], columns = 'count')
      ct = ct/data2.shape[0]
      print(ct)
      ct.plot(kind = 'bar')
```

col_0	count
Pclass	
1	0.242424
2	0.206510
3	0.551066

```
[23]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1096789910>
```



We can see that majority were of class 3

We can see there are 3 unique values for passenger class but these can be converted into categorical feature instead as they refer to a passenger class and also they also have an order to them

```
[24]: Pclass = pd.Categorical(data2["Pclass"],  
                             ordered=True)  
Pclass.describe()
```

```
[24]:
```

	counts	freqs
categories		
1	216	0.242424
2	184	0.206510
3	491	0.551066

```
[25]: data3 = data2.copy()  
data3['Pclass'] = Pclass  
data3.head()
```

```
[25]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked
0	0	3	male	22.0	1	0	7.2500	NaN	S
1	1	1	female	38.0	1	0	71.2833	C85	C
2	1	3	female	26.0	0	0	7.9250	NaN	S
3	1	1	female	35.0	1	0	53.1000	C123	S

4	0	3	male	35.0	0	0	8.0500	NaN	S
---	---	---	------	------	---	---	--------	-----	---

```
[26]: # Do the one hot encoding
data3 = pd.get_dummies(data3, columns=['Pclass'])
data3.head()
```

```
[26]:
```

	Survived	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked	Pclass_1	\
0	0	male	22.0	1	0	7.2500	NaN	S	0	
1	1	female	38.0	1	0	71.2833	C85	C	1	
2	1	female	26.0	0	0	7.9250	NaN	S	0	
3	1	female	35.0	1	0	53.1000	C123	S	1	
4	0	male	35.0	0	0	8.0500	NaN	S	0	

	Pclass_2	Pclass_3
0	0	1
1	0	0
2	0	1
3	0	0
4	0	1

Lets see the survived ratio among the passenger classes

```
[27]: survived_pclass = pd.crosstab(index=data2["Survived"],
                                     columns=data2["Pclass"], margins = True)
display(survived_pclass/survived_pclass.loc["All", "All"])
print(survived_pclass/survived_pclass.loc["All"])
survived_pclass = survived_pclass.div(survived_pclass["All"],
                                       axis=0)
print(survived_pclass)

survived_pclass.pop('All')
survived_pclass.plot(kind = 'bar')
```


Pclass	1	2	3	All
Survived				
0	0.089787	0.108866	0.417508	0.616162
1	0.152637	0.097643	0.133558	0.383838
All	0.242424	0.206510	0.551066	1.000000

Pclass	1	2	3	All
Survived				
0	0.37037	0.527174	0.757637	0.616162
1	0.62963	0.472826	0.242363	0.383838
All	1.00000	1.000000	1.000000	1.000000

Pclass	1	2	3	All
Survived				
0	0.145719	0.176685	0.677596	1.0

```
1          0.397661  0.254386  0.347953  1.0
All        0.242424  0.206510  0.551066  1.0
```

```
[27]: <matplotlib.axes._subplots.AxesSubplot at 0x7f109677eb10>
```



```
[27]:
```

We can see that class 3 people have died a lot when compared to other classes and class 1 people have survived the most when compared to other classes.


3.1.3 Age column

```
[28]: data3['Age'].isnull().sum()
```

```
[28]: 177
```

```
[29]: data3['Age'].plot(kind = 'hist')
```

```
[29]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10961fbf90>
```



as age follows a normal distribution we can switch out the nan values with mean

```
[30]: mean = np.nanmean(data3['Age'])  
mean
```

```
[30]: 29.69911764705882
```


```
[31]: data3['Age'].fillna(mean, inplace = True)
```

```
[32]: data3['Age'].isnull().sum()
```

```
[32]: 0
```

```
[33]: data3['Age'].plot(kind = 'hist')
```

```
[33]: <matplotlib.axes._subplots.AxesSubplot at 0x7f109611f6d0>
```



```
[34]: data3.head()
```

```
[34]:   Survived   Sex  Age  SibSp  Parch    Fare  Cabin Embarked  Pclass_1  \
0         0  male  22.0     1     0   7.2500   NaN        S         0
1         1 female  38.0     1     0  71.2833   C85        C         1
2         1 female  26.0     0     0   7.9250   NaN        S         0
3         1 female  35.0     1     0  53.1000  C123        S         1
4         0  male  35.0     0     0   8.0500   NaN        S         0

      Pclass_2  Pclass_3
0           0         1
1           0         0
2           0         1
3           0         0
4           0         1
```

Using the original data lets what age range survived the most

```
[35]: dummydata = data2.dropna(subset = ["Age"])
dummydata
```


```
[35]:   Survived  Pclass   Sex  Age  SibSp  Parch    Fare  Cabin Embarked
0         0      3  male  22.0     1     0   7.2500   NaN        S
1         1      1 female  38.0     1     0  71.2833   C85        C
2         1      3 female  26.0     0     0   7.9250   NaN        S
3         1      1 female  35.0     1     0  53.1000  C123        S
```

4	0	3	male	35.0	0	0	8.0500	NaN	S
..
885	0	3	female	39.0	0	5	29.1250	NaN	Q
886	0	2	male	27.0	0	0	13.0000	NaN	S
887	1	1	female	19.0	0	0	30.0000	B42	S
889	1	1	male	26.0	0	0	30.0000	C148	C
890	0	3	male	32.0	0	0	7.7500	NaN	Q

[714 rows x 9 columns]

```
[36]: dummydata.boxplot(column = 'Age')
```

```
[36]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10960aba10>
```



```
[37]: bins = [0,20,40,60,80,100,120]
labels = ['0-20','20-40','40-60','60-80','80-100','100+']

dummydata['Age_range'] = pd.cut(dummydata['Age'],bins = bins, labels = labels)
dummydata.head()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:4:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
after removing the cwd from sys.path.

```
[37]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked	\
0	0	3	male	22.0	1	0	7.2500	NaN	S	
1	1	1	female	38.0	1	0	71.2833	C85	C	
2	1	3	female	26.0	0	0	7.9250	NaN	S	
3	1	1	female	35.0	1	0	53.1000	C123	S	
4	0	3	male	35.0	0	0	8.0500	NaN	S	

```
Age_range
0    20-40
1    20-40
2    20-40
3    20-40
4    20-40
```

```
[38]: dummydata['Age_range'].unique()
```

```
[38]: ['20-40', '40-60', '0-20', '60-80']
Categories (6, object): ['0-20' < '20-40' < '40-60' < '60-80' < '80-100' < '100+']
```

```
[39]: survived_age = pd.crosstab(index=dummydata["Survived"],
                                columns=dummydata["Age_range"], margins = True)
display(survived_age/survived_age.loc["All", "All"])
display(survived_age/survived_age.loc["All"])
survived_age = survived_age.div(survived_age["All"],
                                axis=0)
display(survived_age)

survived_age.pop('All')
survived_age.plot(kind = 'bar')
```


Age_range	0-20	20-40	40-60	60-80	All
Survived					
0	0.135854	0.324930	0.109244	0.023810	0.593838
1	0.114846	0.214286	0.070028	0.007003	0.406162
All	0.250700	0.539216	0.179272	0.030812	1.000000

Age_range	0-20	20-40	40-60	60-80	All
Survived					
0	0.541899	0.602597	0.609375	0.772727	0.593838
1	0.458101	0.397403	0.390625	0.227273	0.406162
All	1.000000	1.000000	1.000000	1.000000	1.000000

Age_range	0-20	20-40	40-60	60-80	All
-----------	------	-------	-------	-------	-----

Survived						
0	0.228774	0.547170	0.183962	0.040094	1.0	
1	0.282759	0.527586	0.172414	0.017241	1.0	
All	0.250700	0.539216	0.179272	0.030812	1.0	

[39]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095fd5f90>



3.1.4 Sibsp and parch column

```
[40]: data3['SibSp'].isnull().sum()
```

[40]: 0

```
[41]: data3['Parch'].isnull().sum()
```


[41]: 0

```
[42]: survived_SibSp = pd.crosstab(index=data3["Survived"],
                                   columns=data3["SibSp"])
display(survived_SibSp)
survived_SibSp.plot(kind = 'bar')
```

SibSp	0	1	2	3	4	5	8
-------	---	---	---	---	---	---	---

Survived	0	1	2	3	4	5	6	7
0	398	97	15	12	15	5	7	
1	210	112	13	4	3	0	0	


[42]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095f62e90>



```
[43]: survived_parch = pd.crosstab(index=data3["Survived"],
                                   columns=data3["Parch"])
display(survived_parch)
survived_parch.plot(kind = 'bar')
```

Parch	0	1	2	3	4	5	6
Survived							
0	445	53	40	2	4	4	1
1	233	65	40	3	0	1	0

[43]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095ee88d0>



```
[44]: def with_family(row):
      if (row['SibSp'] !=0) or (row['Parch'] != 0):
          return 1
      else:
          return 0
```

```
[45]: data4 = data3.copy()
      data4['with_family'] = data3.apply(lambda row : with_family(row), axis = 1)
      data4.head()
```

```
[45]:   Survived   Sex   Age  SibSp  Parch    Fare  Cabin Embarked  Pclass_1  \
0         0  male  22.0     1     0   7.2500   NaN      S         0
1         1 female  38.0     1     0  71.2833   C85      C         1
2         1 female  26.0     0     0   7.9250   NaN      S         0
3         1 female  35.0     1     0  53.1000  C123      S         1
4         0  male  35.0     0     0   8.0500   NaN      S         0

      Pclass_2  Pclass_3  with_family
0         0         1             1
1         0         0             1
2         0         1             0
3         0         0             1
4         0         1             0
```


```
[46]: survived_fam = pd.crosstab(index=data4["Survived"],
                                columns=data4["with_family"], margins = True)
display(survived_fam/survived_fam.loc["All", "All"])
display(survived_fam/survived_fam.loc["All"])
survived_fam = survived_fam.div(survived_fam["All"],
                                axis=0)
display(survived_fam)
survived_fam.pop('All')
survived_fam.plot(kind = 'bar')
```

with_family	0	1	All
Survived			
0	0.419753	0.196409	0.616162
1	0.182941	0.200898	0.383838
All	0.602694	0.397306	1.000000

with_family	0	1	All
Survived			
0	0.696462	0.49435	0.616162
1	0.303538	0.50565	0.383838
All	1.000000	1.00000	1.000000

with_family	0	1	All
Survived			
0	0.681239	0.318761	1.0
1	0.476608	0.523392	1.0
All	0.602694	0.397306	1.0

```
[46]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10960333d0>
```




3.1.5 fare column

```
[47]: data3['Fare'].isnull().sum()
```

```
[47]: 0
```


```
[48]: data3['Fare'].plot(kind = 'hist')
```

```
[48]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095dbc0d0>
```



```
[49]: data3.boxplot(column = ['Fare'])
```

```
[49]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095ccd790>
```



```
[50]: index = np.where(data4["Fare"] >100)

data4.loc[index]
```

```
[50]:
```

	Survived	Sex	Age	SibSp	Parch	Fare	Cabin	\
27	0	male	19.000000	3	2	263.0000	C23 C25 C27	
31	1	female	29.699118	1	0	146.5208	B78	
88	1	female	23.000000	3	2	263.0000	C23 C25 C27	
118	0	male	24.000000	0	1	247.5208	B58 B60	
195	1	female	58.000000	0	0	146.5208	B80	
215	1	female	31.000000	1	0	113.2750	D36	
258	1	female	35.000000	0	0	512.3292	NaN	
268	1	female	58.000000	0	1	153.4625	C125	
269	1	female	35.000000	0	0	135.6333	C99	
297	0	female	2.000000	1	2	151.5500	C22 C26	
299	1	female	50.000000	0	1	247.5208	B58 B60	
305	1	male	0.920000	1	2	151.5500	C22 C26	
306	1	female	29.699118	0	0	110.8833	NaN	
307	1	female	17.000000	1	0	108.9000	C65	
311	1	female	18.000000	2	2	262.3750	B57 B59 B63 B66	
318	1	female	31.000000	0	2	164.8667	C7	
319	1	female	40.000000	1	1	134.5000	E34	
325	1	female	36.000000	0	0	135.6333	C32	
332	0	male	38.000000	0	1	153.4625	C91	
334	1	female	29.699118	1	0	133.6500	NaN	
337	1	female	41.000000	0	0	134.5000	E40	
341	1	female	24.000000	3	2	263.0000	C23 C25 C27	
373	0	male	22.000000	0	0	135.6333	NaN	
377	0	male	27.000000	0	2	211.5000	C82	
380	1	female	42.000000	0	0	227.5250	NaN	
390	1	male	36.000000	1	2	120.0000	B96 B98	
393	1	female	23.000000	1	0	113.2750	D36	
435	1	female	14.000000	1	2	120.0000	B96 B98	
438	0	male	64.000000	1	4	263.0000	C23 C25 C27	
498	0	female	25.000000	1	2	151.5500	C22 C26	
505	0	male	18.000000	1	0	108.9000	C65	
527	0	male	29.699118	0	0	221.7792	C95	
537	1	female	30.000000	0	0	106.4250	NaN	
544	0	male	50.000000	1	0	106.4250	C86	
550	1	male	17.000000	0	2	110.8833	C70	
557	0	male	29.699118	0	0	227.5250	NaN	
581	1	female	39.000000	1	1	110.8833	C68	
609	1	female	40.000000	0	0	153.4625	C125	
659	0	male	58.000000	0	2	113.2750	D48	
660	1	male	50.000000	2	0	133.6500	NaN	
679	1	male	36.000000	0	1	512.3292	B51 B53 B55	
689	1	female	15.000000	0	1	211.3375	B5	

698	0	male	49.000000	1	1	110.8833		C68
700	1	female	18.000000	1	0	227.5250		C62 C64
708	1	female	22.000000	0	0	151.5500		NaN
716	1	female	38.000000	0	0	227.5250		C45
730	1	female	29.000000	0	0	211.3375		B5
737	1	male	35.000000	0	0	512.3292		B101
742	1	female	21.000000	2	2	262.3750	B57 B59	B63 B66
763	1	female	36.000000	1	2	120.0000		B96 B98
779	1	female	43.000000	0	1	211.3375		B3
802	1	male	11.000000	1	2	120.0000		B96 B98
856	1	female	45.000000	1	1	164.8667		NaN

	Embarked	Pclass_1	Pclass_2	Pclass_3	with_family
27	S	1	0	0	1
31	C	1	0	0	1
88	S	1	0	0	1
118	C	1	0	0	1
195	C	1	0	0	0
215	C	1	0	0	1
258	C	1	0	0	0
268	S	1	0	0	1
269	S	1	0	0	0
297	S	1	0	0	1
299	C	1	0	0	1
305	S	1	0	0	1
306	C	1	0	0	0
307	C	1	0	0	1
311	C	1	0	0	1
318	S	1	0	0	1
319	C	1	0	0	1
325	C	1	0	0	0
332	S	1	0	0	1
334	S	1	0	0	1
337	C	1	0	0	0
341	S	1	0	0	1
373	C	1	0	0	0
377	C	1	0	0	1
380	C	1	0	0	0
390	S	1	0	0	1
393	C	1	0	0	1
435	S	1	0	0	1
438	S	1	0	0	1
498	S	1	0	0	1
505	C	1	0	0	1
527	S	1	0	0	0
537	C	1	0	0	0
544	C	1	0	0	1

550	C	1	0	0	1
557	C	1	0	0	0
581	C	1	0	0	1
609	S	1	0	0	0
659	C	1	0	0	1
660	S	1	0	0	1
679	C	1	0	0	1
689	S	1	0	0	1
698	C	1	0	0	1
700	C	1	0	0	1
708	S	1	0	0	0
716	C	1	0	0	0
730	S	1	0	0	0
737	C	1	0	0	0
742	C	1	0	0	1
763	S	1	0	0	1
779	S	1	0	0	1
802	S	1	0	0	1
856	S	1	0	0	1

we can see that extra money was payed for passenger class 1 .This outliers gives us this info so we donot the change the outlier

```
[51]: data4.describe()
```

```
[51]:
```

	Survived	Age	SibSp	Parch	Fare	Pclass_1 \
count	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000
mean	0.383838	29.699118	0.523008	0.381594	32.204208	0.242424
std	0.486592	13.002015	1.102743	0.806057	49.693429	0.428790
min	0.000000	0.420000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	22.000000	0.000000	0.000000	7.910400	0.000000
50%	0.000000	29.699118	0.000000	0.000000	14.454200	0.000000
75%	1.000000	35.000000	1.000000	0.000000	31.000000	0.000000
max	1.000000	80.000000	8.000000	6.000000	512.329200	1.000000

	Pclass_2	Pclass_3	with_family
count	891.000000	891.000000	891.000000
mean	0.206510	0.551066	0.397306
std	0.405028	0.497665	0.489615
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	1.000000	0.000000
75%	0.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000

```
[52]: bins = [0,20,40,60,80,100,520]
labels = ['0-20','20-40','40-60','60-80','80-100','100-520']
```

```
fare_range = pd.cut(data4['Fare'],bins = bins, labels = labels)
```


```
[53]: survived_fare = pd.crosstab(index=data4["Survived"],
                                columns=fare_range, margins = True)
display(survived_fare/survived_fare.loc["All","All"])
display(survived_fare/survived_fare.loc["All"])
survived_fare = survived_fare.div(survived_fare["All"],
                                axis=0)
display(survived_fare)
survived_fare.pop('All')
survived_fare.plot(kind = 'bar')
```

Fare	0-20	20-40	40-60	60-80	80-100	100-520	All
Survived							
0	0.408676	0.130137	0.026256	0.026256	0.003425	0.015982	0.610731
1	0.162100	0.098174	0.035388	0.028539	0.020548	0.044521	0.389269
All	0.570776	0.228311	0.061644	0.054795	0.023973	0.060502	1.000000

Fare	0-20	20-40	40-60	60-80	80-100	100-520	All
Survived							
0	0.716	0.57	0.425926	0.479167	0.142857	0.264151	0.610731
1	0.284	0.43	0.574074	0.520833	0.857143	0.735849	0.389269
All	1.000	1.00	1.000000	1.000000	1.000000	1.000000	1.000000

Fare	0-20	20-40	40-60	60-80	80-100	100-520	All
Survived							
0	0.669159	0.213084	0.042991	0.042991	0.005607	0.026168	1.0
1	0.416422	0.252199	0.090909	0.073314	0.052786	0.114370	1.0
All	0.570776	0.228311	0.061644	0.054795	0.023973	0.060502	1.0

```
[53]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095cc1e10>
```



3.2 categorical columns

```
[54]: categorical_cols = data2.dtypes[data.dtypes == "object"].index
      print(categorical_cols)
```

```
Index(['Sex', 'Cabin', 'Embarked'], dtype='object')
```

3.2.1 Sex column

```
[55]: data4['Sex'].unique()
```

```
[55]: array(['male', 'female'], dtype=object)
```

```
[56]: data4['Sex'].isnull().sum()
```

```
[56]: 0
```

```
[57]: survivor_sex = pd.crosstab(index=data4["Survived"],
                                columns=data["Sex"], margins = True)
      display(survivor_sex/survivor_sex.loc["All", "All"])
      display(survivor_sex/survivor_sex.loc["All"])
      survivor_sex = survivor_sex.div(survivor_sex["All"],
```

```

        axis=0)
display(surviver_sex)
surviver_sex.pop('All')
surviver_sex.plot(kind = 'bar')


```

Sex	female	male	All
Survived			
0	0.090909	0.525253	0.616162
1	0.261504	0.122334	0.383838
All	0.352413	0.647587	1.000000

Sex	female	male	All
Survived			
0	0.257962	0.811092	0.616162
1	0.742038	0.188908	0.383838
All	1.000000	1.000000	1.000000

Sex	female	male	All
Survived			
0	0.147541	0.852459	1.0
1	0.681287	0.318713	1.0
All	0.352413	0.647587	1.0

[57]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095b89090>



```
[58]: # Do the one hot encoding
data4 = pd.get_dummies(data4, columns=['Sex'])
data4.head()
```

```
[58]:
```

	Survived	Age	SibSp	Parch	Fare	Cabin	Embarked	Pclass_1	Pclass_2	\
0	0	22.0	1	0	7.2500	NaN	S	0	0	
1	1	38.0	1	0	71.2833	C85	C	1	0	
2	1	26.0	0	0	7.9250	NaN	S	0	0	
3	1	35.0	1	0	53.1000	C123	S	1	0	
4	0	35.0	0	0	8.0500	NaN	S	0	0	

	Pclass_3	with_family	Sex_female	Sex_male
0	1	1	0	1
1	0	1	1	0
2	1	0	1	0
3	0	1	1	0
4	1	0	0	1

3.2.2 Cabin

```
[59]: data4['Cabin'].isnull().sum()
```

```
[59]: 687
```

```
[60]: data4['Cabin'].unique()
```

```
[60]: array([nan, 'C85', 'C123', 'E46', 'G6', 'C103', 'D56', 'A6',
        'C23 C25 C27', 'B78', 'D33', 'B30', 'C52', 'B28', 'C83', 'F33',
        'F G73', 'E31', 'A5', 'D10 D12', 'D26', 'C110', 'B58 B60', 'E101',
        'F E69', 'D47', 'B86', 'F2', 'C2', 'E33', 'B19', 'A7', 'C49', 'F4',
        'A32', 'B4', 'B80', 'A31', 'D36', 'D15', 'C93', 'C78', 'D35',
        'C87', 'B77', 'E67', 'B94', 'C125', 'C99', 'C118', 'D7', 'A19',
        'B49', 'D', 'C22 C26', 'C106', 'C65', 'E36', 'C54',
        'B57 B59 B63 B66', 'C7', 'E34', 'C32', 'B18', 'C124', 'C91', 'E40',
        'T', 'C128', 'D37', 'B35', 'E50', 'C82', 'B96 B98', 'E10', 'E44',
        'A34', 'C104', 'C111', 'C92', 'E38', 'D21', 'E12', 'E63', 'A14',
        'B37', 'C30', 'D20', 'B79', 'E25', 'D46', 'B73', 'C95', 'B38',
        'B39', 'B22', 'C86', 'C70', 'A16', 'C101', 'C68', 'A10', 'E68',
        'B41', 'A20', 'D19', 'D50', 'D9', 'A23', 'B50', 'A26', 'D48',
        'E58', 'C126', 'B71', 'B51 B53 B55', 'D49', 'B5', 'B20', 'F G63',
        'C62 C64', 'E24', 'C90', 'C45', 'E8', 'B101', 'D45', 'C46', 'D30',
        'E121', 'D11', 'E77', 'F38', 'B3', 'D6', 'B82 B84', 'D17', 'A36',
        'B102', 'B69', 'E49', 'C47', 'D28', 'E17', 'A24', 'C50', 'B42',
        'C148'], dtype=object)
```

even though this has a lot of nan values we cant delete this columns because dependind on which deck got flooded first the survival will on that so we assign a new cabin called a No cabin to this column for nan and also there seems to be repeating cabin letters which are unique so we can filter out the numbers

```
[61]: cabin = data4["Cabin"].astype(str)

Cabin = np.array([obs[0] for obs in cabin])
ct = pd.crosstab(index = Cabin, columns = 'count')
display(ct)
Cabin = pd.Categorical(Cabin)

data5 = data4.copy()
data5['Cabin'] = Cabin
data5 = pd.get_dummies(data5, columns=['Cabin'])
data5.head()
```

```
col_0  count
row_0
A      15
B      47
C      59
D      33
E      32
F      13
G       4
T       1
n     687
```

```
[61]:
```

	Survived	Age	SibSp	Parch	Fare	Embarked	Pclass_1	Pclass_2	\
0	0	22.0	1	0	7.2500	S	0	0	
1	1	38.0	1	0	71.2833	C	1	0	
2	1	26.0	0	0	7.9250	S	0	0	
3	1	35.0	1	0	53.1000	S	1	0	
4	0	35.0	0	0	8.0500	S	0	0	

	Pclass_3	with_family	...	Sex_male	Cabin_A	Cabin_B	Cabin_C	Cabin_D	\
0	1	1	...	1	0	0	0	0	
1	0	1	...	0	0	0	1	0	
2	1	0	...	0	0	0	0	0	
3	0	1	...	0	0	0	1	0	
4	1	0	...	1	0	0	0	0	

	Cabin_E	Cabin_F	Cabin_G	Cabin_T	Cabin_n
0	0	0	0	0	1
1	0	0	0	0	0
2	0	0	0	0	1

3	0	0	0	0	0
4	0	0	0	0	1

[5 rows x 21 columns]

```
[62]: survivor_cabin = pd.crosstab(index=data4["Survived"],
                                   columns=Cabin, margins = True)
display(survivor_cabin/survivor_cabin.loc["All", "All"])
display(survivor_cabin/survivor_cabin.loc["All"])
survivor_cabin = survivor_cabin.div(survivor_cabin["All"],
                                   axis=0)
display(survivor_cabin)
survivor_cabin.pop('All')
survivor_cabin.plot(kind = 'bar')
```

col_0	A	B	C	D	E	F \
Survived						
0	0.008979	0.013468	0.026936	0.008979	0.008979	0.005612
1	0.007856	0.039282	0.039282	0.028058	0.026936	0.008979
All	0.016835	0.052750	0.066218	0.037037	0.035915	0.014590

col_0	G	T	n	All
Survived				
0	0.002245	0.001122	0.539843	0.616162
1	0.002245	0.000000	0.231201	0.383838
All	0.004489	0.001122	0.771044	1.000000

col_0	A	B	C	D	E	F	G	T \
Survived								
0	0.533333	0.255319	0.40678	0.242424	0.25	0.384615	0.5	1.0
1	0.466667	0.744681	0.59322	0.757576	0.75	0.615385	0.5	0.0
All	1.000000	1.000000	1.00000	1.000000	1.00	1.000000	1.0	1.0


col_0	n	All
Survived		
0	0.700146	0.616162
1	0.299854	0.383838
All	1.000000	1.000000

col_0	A	B	C	D	E	F \
Survived						
0	0.014572	0.021858	0.043716	0.014572	0.014572	0.009107
1	0.020468	0.102339	0.102339	0.073099	0.070175	0.023392
All	0.016835	0.052750	0.066218	0.037037	0.035915	0.014590

col_0	G	T	n	All
-------	---	---	---	-----

Survived					
0	0.003643	0.001821	0.876138	1.0	
1	0.005848	0.000000	0.602339	1.0	
All	0.004489	0.001122	0.771044	1.0	

[62]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095be4a50>



3.2.3 embark column

[63]: data5['Embarked'].unique()

[63]: array(['S', 'C', 'Q', nan], dtype=object)

[64]: data5['Embarked'].isnull().sum()

[64]: 2

```
[65]: survivor_embarked = pd.crosstab(index=data4["Survived"],
                                     columns=data["Embarked"], margins = True)
display(survivor_embarked/survivor_embarked.loc["All","All"])
display(survivor_embarked/survivor_embarked.loc["All"])
survivor_embarked = survivor_embarked.div(survivor_embarked["All"],
                                     axis=0)
```




```
display(surviver_embarked)
surviver_embarked.pop('All')
surviver_embarked.plot(kind = 'bar')
```

Embarked	C	Q	S	All
Survived				
0	0.084364	0.052868	0.480315	0.617548
1	0.104612	0.033746	0.244094	0.382452
All	0.188976	0.086614	0.724409	1.000000

Embarked	C	Q	S	All
Survived				
0	0.446429	0.61039	0.663043	0.617548
1	0.553571	0.38961	0.336957	0.382452
All	1.000000	1.00000	1.000000	1.000000

Embarked	C	Q	S	All
Survived				
0	0.136612	0.085610	0.777778	1.0
1	0.273529	0.088235	0.638235	1.0
All	0.188976	0.086614	0.724409	1.0

[65]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095a0ae90>



```
[66]: # Do the one hot encoding
data6 = data5.copy()
data6 = pd.get_dummies(data6, columns=['Embarked'])
data6.head()
```

```
[66]:
```

	Survived	Age	SibSp	Parch	Fare	Pclass_1	Pclass_2	Pclass_3	\
0	0	22.0	1	0	7.2500	0	0	1	
1	1	38.0	1	0	71.2833	1	0	0	
2	1	26.0	0	0	7.9250	0	0	1	
3	1	35.0	1	0	53.1000	1	0	0	
4	0	35.0	0	0	8.0500	0	0	1	


	with_family	Sex_female	...	Cabin_C	Cabin_D	Cabin_E	Cabin_F	Cabin_G	\
0	1	0	...	0	0	0	0	0	
1	1	1	...	1	0	0	0	0	
2	0	1	...	0	0	0	0	0	
3	1	1	...	1	0	0	0	0	
4	0	0	...	0	0	0	0	0	

	Cabin_T	Cabin_n	Embarked_C	Embarked_Q	Embarked_S
0	0	1	0	0	1
1	0	0	1	0	0
2	0	1	0	0	1
3	0	0	0	0	1
4	0	1	0	0	1

[5 rows x 23 columns]

```
[67]: import matplotlib.pyplot as plt
fig, ax = plt.subplots(figsize=(8,8))
corr = data.corr()
sns.heatmap(corr, annot=True, ax = ax)
```

```
[67]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1095f1c7d0>
```



3.3 Hypothesis testing


3.3.1 H0: Data is normally distributed H1: Data is not normally distributed

```
[68]: import scipy.stats as ss
```

```
[69]: dummy = data[data['Age'].notna()].Age
```

```
[70]: dummy.plot(kind = 'hist')
```

```
[70]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1093047a10>
```



```
[71]: from scipy.stats import normaltest
stat, pvalue = normaltest(dummy)
print(stat, pvalue)
if pvalue<0.05:
    print("not a normal distribution")
else:
    print("Normal distribution")
```

```
18.105032952089758 0.00011709599657350757
not a normal distribution
```


We apply boxcox to make the data normally distributed

```
[72]: from scipy.stats import boxcox
```

```
[73]: normal_Age = boxcox(dummy)
```

```
[74]: pd.DataFrame(normal_Age[0]).plot(kind = 'hist')
```

```
[74]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1092ff2410>
```



```
[75]: from scipy.stats import normaltest
stat, pvalue = normaltest(normal_Age[0])
print(stat, pvalue)
if pvalue<0.05:
    print("not a normal distribution")
else:
    print("Normal distribution")
```

```
1.2384147036970252 0.5383710072962454
Normal distribution
```

3.3.2 H0: Sex doesn't matter for survival H1: Sex matters for survival

```
[76]: survivor_sex = pd.crosstab(index=data4["Survived"],
                                columns=data['Sex'])
```

```
[77]: survivor_sex
```

```
[77]: Sex      female  male
Survived
0         81    468
1        233    109
```

```
[78]: from scipy.stats import chi2_contingency
stat, p, dof, expected = chi2_contingency(survivor_sex)
```

```

print(p)
if p < 0.05:
    print('Sex matters in terms of survival')
else:
    print('Sex doesnt matter for survival')

```

```

1.1973570627755645e-58
Sex matters in terms of survival

```

3.3.3 H0:there is no relationship between age and fare H1: there is a relationship between age and fare

```

[79]: from scipy.stats import spearmanr
coef, p = spearmanr(data6[1:60]['Age'], data6[1:60]['Fare'])
print(p)
if p<0.05:
    print('there is a relationship between age and Fare')
else:
    print('there is no relationship between age and fare')

```

```

0.47931972164440195
there is no relationship between age and fare

```

4 Data splitting

```

[80]: feature_cols = list(data6.columns)
feature_cols.remove('Survived')
strat_shuf_split = StratifiedShuffleSplit(n_splits=1,
                                         test_size=0.3,
                                         random_state=42)

train_index, test_index = next(strat_shuf_split.split(data6[feature_cols],
→data6.Survived))

# Create the dataframes
X_train = data6.loc[train_index, feature_cols]
y_train = data6.loc[train_index, 'Survived']

X_test = data6.loc[test_index, feature_cols]
y_test = data6.loc[test_index, 'Survived']
len(X_test), len(X_train)

```

```

[80]: (268, 623)

```

```
[81]: X = data6[feature_cols]
      y = data6['Survived']
```

5 Logistic Regression

Choosing best scaling method

```
[164]: from sklearn.preprocessing import StandardScaler, MinMaxScaler, \
      ↪RobustScaler, PolynomialFeatures
scalers = {'standards': StandardScaler(),
          'minmax': MinMaxScaler(),
          'robust': RobustScaler()}
LR = LogisticRegression()
scores = {}
for scaler_label, scaler in scalers.items():
    trainingset = scaler.fit_transform(X_train)
    testset = scaler.transform(X_test)
    LR.fit(trainingset, y_train)
    predictions = LR.predict(testset)
    key = scaler_label + 'Scaling'
    scores[key] = roc_auc_score(y_test, predictions)

for key, val in scores.items():
    print(key, val)
```

standardsScaling 0.7901441600470727

minmaxScaling 0.7822594880847309

robustScaling 0.7931744630773758

robustScaling works best for logistic regression model

Hyperparameter tuning for polynomial features

```
[87]: from sklearn.model_selection import GridSearchCV
      from sklearn.model_selection import KFold, cross_val_predict
      from sklearn.pipeline import Pipeline
sf = StratifiedKFold(n_splits=3, shuffle = True, random_state=72018)
estimator = Pipeline([("Polynomial", PolynomialFeatures()),
                      ("RobustScaler", RobustScaler()),
                      ("LogisticRegression", LogisticRegression(max_iter = \
      ↪10000000))])
params = {"Polynomial__degree": np.arange(1,5,1),
          "LogisticRegression__solver": ["liblinear"]}
grid = GridSearchCV(estimator, params, cv =sf, scoring = 'accuracy' )
grid.fit(X, y)
print(grid.best_params_, grid.best_score_)
```

```
{'LogisticRegression__solver': 'liblinear', 'Polynomial__degree': 1}  
0.8058361391694725
```

hyperparameter tuning of logistic regression

```
[88]: params = {"LogisticRegression__C": np.logspace(-4, 4, 20),  
               "LogisticRegression__penalty": ['l1', 'l2'],  
               "LogisticRegression__solver": ["liblinear"]}  
estimator = Pipeline([("RobustScaler", RobustScaler()),  
                       ("LogisticRegression", LogisticRegression())])  
grid = GridSearchCV(estimator, params, cv=sf, scoring = 'accuracy')  
grid.fit(X, y)  
print(grid.best_params_, grid.best_score_)  
C = grid.best_params_['LogisticRegression__C']  
penalty = grid.best_params_['LogisticRegression__penalty']  
solver = grid.best_params_['LogisticRegression__solver']  
lr = LogisticRegression(C = C, penalty = penalty, solver = solver)  
lr.fit(X_train, y_train)
```

```
{'LogisticRegression__C': 0.615848211066026, 'LogisticRegression__penalty':  
'l2', 'LogisticRegression__solver': 'liblinear'} 0.809203142536476
```

```
[88]: LogisticRegression(C=0.615848211066026, solver='liblinear')
```

Model evaluation

```
[156]: pred_lr = lr.predict(X_test)  
metrics = pd.DataFrame()  
precision_lr, recall_lr, f1score, _ = score(y_test, pred_lr, average='weighted')  
# adding lr stats to metrics DataFrame  
lr_stats = pd.Series({'precision': precision_lr,  
                     'recall': recall_lr,  
                     'accuracy': accuracy_score(y_test, pred_lr),  
                     'f1score': f1_score(y_test, pred_lr, average='weighted'),  
                     'auc': roc_auc_score(y_test, pred_lr, average='weighted')},  
                     name='Logistic Regression')  
  
# Report outcomes  
display(lr_stats)  
print(classification_report(y_test, pred_lr))
```

```
precision    0.815939  
recall       0.817164  
accuracy     0.817164  
f1score      0.816255  
auc          0.802265  
Name: Logistic Regression, dtype: float64
```

```
precision    recall  f1-score   support
```


0	0.84	0.87	0.85	165
1	0.78	0.74	0.76	103
accuracy			0.82	268
macro avg	0.81	0.80	0.80	268
weighted avg	0.82	0.82	0.82	268

```
[116]: score_lr_train=np.mean(cross_val_score(tree, X_train, y_train, cv=sf,
    ↳scoring='accuracy'))
print("Training Accuracy is : ",score_lr_train)
score_lr_test=np.mean(cross_val_score(tree, X_test, y_test, cv=sf,
    ↳scoring='accuracy'))
print("Testing Accuracy is : ",score_lr_test)
```

Training Accuracy is : 0.7881518642388207
 Testing Accuracy is : 0.7984602580108199

6 KNN classifier

Scaling Selection

```
[90]: from sklearn.preprocessing import StandardScaler, MinMaxScaler,
    ↳RobustScaler,PolynomialFeatures
scalers = {'standars': StandardScaler(),
    'minmax': MinMaxScaler(),
    'robust':RobustScaler()}
LR = KNeighborsClassifier()
scores = {}
for scaler_label, scaler in scalers.items():
    trainingset = scaler.fit_transform(X_train)
    testset = scaler.transform(X_test)
    LR.fit(trainingset, y_train)
    predictions = LR.predict(testset)
    key = scaler_label + 'Scaling'
    scores[key] = roc_auc_score(y_test, predictions)

for key, val in scores.items():
    print(key, val)
```

standarsScaling 0.7616357752280083
 minmaxScaling 0.7537216828478965
 robustScaling 0.7337157987643426

hyperparameter tuning for polynomial features

```
[91]: from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import KFold, cross_val_predict
from sklearn.pipeline import Pipeline
sf = StratifiedKFold(n_splits=3, shuffle = True, random_state=72018)
estimator = Pipeline([("Polynomial", PolynomialFeatures()),
                      ("StandardScaler", StandardScaler()),
                      ("LogisticRegression", KNeighborsClassifier())])
params = {"Polynomial__degree": np.arange(1, 5, 1)}
grid = GridSearchCV(estimator, params, cv=sf, scoring = 'accuracy' )
grid.fit(X, y)
print(grid.best_params_, grid.best_score_)
```

```
{'Polynomial__degree': 1} 0.7845117845117845
```

hyperparameter tuning of KNN

```
[94]: params = {"KNN__n_neighbors": np.arange(1, 30, 1)}
estimator = Pipeline([("StandardScaler", StandardScaler()),
                      ("KNN", KNeighborsClassifier())])
grid = GridSearchCV(estimator, params, cv=sf, scoring = 'accuracy')
grid.fit(X, y)
print(grid.best_params_, grid.best_score_)
C = grid.best_params_['KNN__n_neighbors']

knn = KNeighborsClassifier(n_neighbors=C)
knn.fit(X_train, y_train)
```

```
{'KNN__n_neighbors': 12} 0.7968574635241302
```

```
[94]: KNeighborsClassifier(n_neighbors=12)
```

Model evaluation

```
[196]: pred_knn = knn.predict(X_test)
precision_lr, recall_lr, f1score, _ = score(y_test, pred_knn, average='weighted')
# adding lr stats to metrics DataFrame
knn_stats = pd.Series({'precision': precision_lr,
                      'recall': recall_lr,
                      'accuracy': accuracy_score(y_test, pred_knn),
                      'f1score': f1_score(y_test, pred_knn, average='weighted'),
                      'auc': roc_auc_score(y_test,
→pred_knn, average='weighted')},
                      name='KNN')

# Report outcomes
display(knn_stats)
print(classification_report(y_test, pred_knn))
```

```
precision    0.674231
recall      0.682836
```

```

accuracy    0.682836
f1score     0.673237
auc         0.643925
Name: KNN, dtype: float64

```

	precision	recall	f1-score	support
0	0.71	0.81	0.76	165
1	0.61	0.48	0.54	103
accuracy			0.68	268
macro avg	0.66	0.64	0.65	268
weighted avg	0.67	0.68	0.67	268

```

[118]: score_knn_train=np.mean(cross_val_score(knn, X_train, y_train, cv=sf,
        ↳scoring='accuracy'))
print("Training Accuracy is : ",score_knn_train)
score_knn_test=np.mean(cross_val_score(knn, X_test, y_test, cv=sf,
        ↳scoring='accuracy'))
print("Testing Accuracy is : ",score_knn_test)

```

```

Training Accuracy is :  0.6950250836120402
Testing Accuracy is :  0.5969621306699958

```

7 SVM

Scaling Selection

```

[96]: from sklearn.preprocessing import StandardScaler, MinMaxScaler,
        ↳RobustScaler,PolynomialFeatures
scalers = {'standards': StandardScaler(),
           'minmax': MinMaxScaler(),
           'robust':RobustScaler()}
LR = SVC()
scores = {}
for scaler_label, scaler in scalers.items():
    trainingset = scaler.fit_transform(X_train)
    testset = scaler.transform(X_test)
    LR.fit(trainingset, y_train)
    predictions = LR.predict(testset)
    key = scaler_label + 'Scaling'
    scores[key] = roc_auc_score(y_test, predictions)

for key, val in scores.items():
    print(key, val)

```

```
standardsScaling 0.7470726684318918
minmaxScaling 0.7719035010297147
robustScaling 0.794380700205943
```

Hyperparameter tuning for polynomial features

```
[97]: from sklearn.model_selection import GridSearchCV
      from sklearn.model_selection import KFold, cross_val_predict
      from sklearn.pipeline import Pipeline
      sf = StratifiedKFold(n_splits=3, shuffle = True, random_state=72018)
      estimator = Pipeline([("Polynomial", PolynomialFeatures()),
                            ("RobustScaler", RobustScaler()),
                            ("svn", SVC())])
      params = {"Polynomial__degree": np.arange(1, 5, 1)}
      grid = GridSearchCV(estimator, params, cv = sf, scoring = 'accuracy' )
      grid.fit(X, y)
      print(grid.best_params_, grid.best_score_)
```

```
{'Polynomial__degree': 1} 0.8103254769921436
```

hyperparameter tuning for svm

```
[98]: estimator = Pipeline([("RobustScaler", RobustScaler()),
                            ("svn", SVC())])
      params = {"svn__C": (0.1, 0.5, 1, 2, 5, 10, 20),
                "svn__gamma": (0.001, 0.01, 0.1, 0.25, 0.5, 0.75, 1),
                "svn__kernel": ('linear', 'poly', 'rbf'),
                }
      grid = GridSearchCV(estimator, params, cv = sf, scoring = 'accuracy' )
      grid.fit(X, y)
      print(grid.best_params_, grid.best_score_)

      C = grid.best_params_['svn__C']
      gamma = grid.best_params_['svn__gamma']
      kernel = grid.best_params_['svn__kernel']

      svm = SVC(C = C, gamma = gamma, kernel = kernel)
      svm.fit(X_train, y_train)
```

```
{'svn__C': 1, 'svn__gamma': 0.25, 'svn__kernel': 'rbf'} 0.8181818181818182
```

```
[98]: SVC(C=1, gamma=0.25)
```

Model evaluation

```
[155]: pred_svm = svm.predict(X_test)
      precision_lr, recall_lr, f1score, _ = score(y_test, pred_svm, average='weighted')
      svm_stats = pd.Series({'precision': precision_lr,
                            'recall': recall_lr,
```

```

        'accuracy': accuracy_score(y_test, pred_svm),
        'f1score': f1_score(y_test, pred_svm, average='weighted'),
        'auc': roc_auc_score(y_test,
→pred_svm, average='weighted')},
        name='Support vector machine')
# Report outcomes
display(svm_stats)
print(classification_report(y_test, pred_svm))

```

```

precision    0.652488
recall       0.664179
accuracy     0.664179
f1score      0.640135
auc          0.606884
Name: Support vector machine, dtype: float64

```

	precision	recall	f1-score	support
0	0.68	0.85	0.76	165
1	0.61	0.36	0.45	103
accuracy			0.66	268
macro avg	0.64	0.61	0.60	268
weighted avg	0.65	0.66	0.64	268

```

[154]: score_svm_train=np.mean(cross_val_score(svm, X_train, y_train, cv=sf,
→scoring='accuracy'))
print("Training Accuracy is : ",score_svm_train)
score_svm_test=np.mean(cross_val_score(svm, X_test, y_test, cv=sf,
→scoring='accuracy'))
print("Testing Accuracy is : ",score_svm_test)

```

```

Training Accuracy is :  0.6709479127957388
Testing Accuracy is :  0.6156471077819393

```

[109]:

```

Testing Accuracy is :  0.6156471077819393

```

8 Decision tree

Scaling selection

```
[86]: from sklearn.preprocessing import StandardScaler, MinMaxScaler, \
      ↪ RobustScaler, PolynomialFeatures
scalers = {'standards': StandardScaler(),
          'minmax': MinMaxScaler(),
          'robust': RobustScaler()}
LR = DecisionTreeClassifier()
scores = {}
for scaler_label, scaler in scalers.items():
    trainingset = scaler.fit_transform(X_train)
    testset = scaler.transform(X_test)
    LR.fit(trainingset, y_train)
    predictions = LR.predict(testset)
    key = scaler_label + 'Scaling'
    scores[key] = roc_auc_score(y_test, predictions)

for key, val in scores.items():
    print(key, val)
```

```
standardsScaling 0.7464548396587232
minmaxScaling 0.7318917328626066
robustScaling 0.7306854957340395
```

Hyperparameter tuning for polynomial features

```
[105]: from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import KFold, cross_val_predict
from sklearn.pipeline import Pipeline
sf = StratifiedKFold(n_splits=3, shuffle = True, random_state=72018)
estimator = Pipeline([("Polynomial", PolynomialFeatures()),
                      ("Scaler", StandardScaler()),
                      ("DT", DecisionTreeClassifier())])
params = {"Polynomial__degree": np.arange(1, 5, 1)}
grid = GridSearchCV(estimator, params, cv = sf, scoring = 'accuracy' )
grid.fit(X, y)
print(grid.best_params_, grid.best_score_)
```

```
{'Polynomial__degree': 2} 0.7789001122334457
```

hyperparameter tuning for decision trees

```
[132]: params = {"dt__criterion": ("gini", "entropy"),
                "dt__splitter": ("best", "random"),
                "dt__max_depth": (list(range(1, 25))),
                "dt__min_samples_split": [2, 3, 4],
                "dt__min_samples_leaf": list(range(1, 20))
                }
estimator = \
    ↪ Pipeline([("Polynomial", PolynomialFeatures(degree=2)), ("Scaler", StandardScaler()),
```

```

        ("dt", DecisionTreeClassifier(random_state=42))])
grid = GridSearchCV(estimator, params, scoring="accuracy",cv=sf)
grid.fit(X, y)
print(grid.best_params_, grid.best_score_)
criterion = grid.best_params_['dt__criterion']
splitter = grid.best_params_['dt__splitter']
max_depth = grid.best_params_['dt__max_depth']
min_samples_split = grid.best_params_['dt__min_samples_split']
min_samples_leaf = grid.best_params_['dt__min_samples_leaf']

```

```

{'dt__criterion': 'entropy', 'dt__max_depth': 7, 'dt__min_samples_leaf': 5,
'dt__min_samples_split': 2, 'dt__splitter': 'random'} 0.8204264870931537

```

```

[133]: tree = DecisionTreeClassifier(criterion = criterion, splitter = splitter,
    ↳max_depth = max_depth, min_samples_split = min_samples_split,
    ↳min_samples_leaf = min_samples_leaf )
tree.fit(X_train, y_train)

```

```

[133]: DecisionTreeClassifier(criterion='entropy', max_depth=7, min_samples_leaf=5,
    splitter='random')

```

```

[150]: pred_dt = tree.predict(X_test)
precision_lr, recall_lr, f1score, _ = score(y_test,pred_dt,average='weighted')
dt_stats = pd.Series({'precision':precision_lr,
    'recall':recall_lr,
    'accuracy':accuracy_score(y_test, pred_dt),
    'f1score':f1_score(y_test, pred_dt,average='weighted'),
    'auc': roc_auc_score(y_test, pred_dt,average='weighted')},
    name='decision Tree')

# Report outcomes
display(dt_stats)
print(classification_report(y_test, pred_dt))

```

```

precision    0.795415
recall       0.794776
accuracy     0.794776
f1score      0.788565
auc          0.762195
Name: KNN, dtype: float64

```

	precision	recall	f1-score	support
0	0.79	0.90	0.84	165
1	0.80	0.62	0.70	103
accuracy			0.79	268
macro avg	0.80	0.76	0.77	268

weighted avg 0.80 0.79 0.79 268

```
[153]: score_tree_train=np.mean(cross_val_score(tree, X_train, y_train, cv=sf,
    ↳scoring='accuracy'))
print("Training Accuracy is : ",score_tree_train)
score_tree_test=np.mean(cross_val_score(tree, X_test, y_test, cv=sf,
    ↳scoring='accuracy'))
print("Testing Accuracy is : ",score_tree_test)
```

Training Accuracy is : 0.7817106404062927
Testing Accuracy is : 0.7908447773616314

9 RandomForest

```
[136]: from sklearn.preprocessing import StandardScaler, MinMaxScaler,
    ↳RobustScaler,PolynomialFeatures
scalers = {'standards': StandardScaler(),
           'minmax': MinMaxScaler(),
           'robust':RobustScaler()}
LR = RandomForestClassifier()
scores = {}
for scaler_label, scaler in scalers.items():
    trainingset = scaler.fit_transform(X_train)
    testset = scaler.transform(X_test)
    LR.fit(trainingset, y_train)
    predictions = LR.predict(testset)
    key = scaler_label + 'Scaling'
    scores[key] = roc_auc_score(y_test, predictions)

for key, val in scores.items():
    print(key, val)
```

standardsScaling 0.7622241835834069
minmaxScaling 0.7701088555457488
robustScaling 0.7670785525154457

```
[142]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

estimator =
    ↳Pipeline([("Polynomial",PolynomialFeatures(degree=2)),("Scaler",MinMaxScaler()),
               ("rf", RandomForestClassifier(max_features='auto',
    ↳oob_score=True, random_state=1, n_jobs=-1))])
params = {"rf__criterion":("gini", "entropy"),
```



```

    "rf__max_depth":(list(range(1, 25))),
    "rf__min_samples_split":[2, 3, 4,10,12,16],
    "rf__min_samples_leaf":list(range(1, 20)),
    "rf__n_estimators": [50, 100, 400, 700, 1000]
}

```

```

grid = GridSearchCV(estimator, params, scoring='accuracy', cv=sf, n_jobs=-1)

grid = grid.fit(X, y)

print( grid.best_score_, grid.best_params_,)

```

```

0.838383838384 {'min_samples_split': 10, 'n_estimators': 700, 'criterion':
'gini', 'min_samples_leaf': 1}

```

```

[144]: rfc = RandomForestClassifier(min_samples_split = 10, n_estimators = 700,
    ↪criterion = 'gini',min_samples_leaf= 1)
rfc.fit(X_train, y_train)

```

```

[144]: RandomForestClassifier(min_samples_split=10, n_estimators=700)

```

```

[151]: pred_rfc = rfc.predict(X_test)
precision_lr, recall_lr, f1score, _ = score(y_test,pred_rfc,average='weighted')

rfc_stats = pd.Series({'precision':precision_lr,
                      'recall':recall_lr,
                      'accuracy':accuracy_score(y_test, pred_rfc),
                      'f1score':f1_score(y_test, pred_rfc,average='weighted'),
                      'auc': roc_auc_score(y_test,
    ↪pred_rfc,average='weighted'))},
                      name='Random forest')

# Report outcomes
display(rfc_stats)
print(classification_report(y_test, pred_rfc))

```

```

precision    0.809013
recall       0.809701
accuracy     0.809701
f1score      0.805805
auc          0.783436
Name: Random forest, dtype: float64

```

	precision	recall	f1-score	support
0	0.81	0.90	0.85	165
1	0.80	0.67	0.73	103

accuracy			0.81	268
macro avg	0.81	0.78	0.79	268
weighted avg	0.81	0.81	0.81	268

```
[152]: score_rfc_train=np.mean(cross_val_score(rfc, X_train, y_train, cv=sf,
→scoring='accuracy'))
print("Training Accuracy is : ",score_rfc_train)
score_rfc_test=np.mean(cross_val_score(rfc, X_test, y_test, cv=sf,
→scoring='accuracy'))
print("Testing Accuracy is : ",score_rfc_test)
```

Training Accuracy is : 0.8009801189149015
Testing Accuracy is : 0.813358302122347

10 Model selection

```
[187]: training = [score_lr_train, score_knn_train, score_svm_train, score_tree_train,
→score_rfc_train]
testing = [score_lr_test, score_knn_test, score_svm_test, score_tree_test,
→score_rfc_test]
```

```
[192]: accuracy = pd.DataFrame([training, testing], columns = ['LogisticRegression', 'K_
→neighbours', 'Support vector machine', 'Decision tree', 'Random forest_
→classifier'], index = ['Training', 'Testing'])
```

```
[163]: metrics = pd.DataFrame()
metrics.append([lr_stats.T, knn_stats.T, svm_stats.T, dt_stats.T, rfc_stats.T])
```

```
[163]:
```

	precision	recall	accuracy	f1score	auc
Logistic Regression	0.815939	0.817164	0.817164	0.816255	0.802265
KNN	0.674231	0.682836	0.682836	0.673237	0.643925
Support vector machine	0.652488	0.664179	0.664179	0.640135	0.606884
KNN	0.795415	0.794776	0.794776	0.788565	0.762195
Random forest	0.809013	0.809701	0.809701	0.805805	0.783436

```
[194]: accuracy.T
```

```
[194]:
```

	Training	Testing
LogisticRegression	0.788152	0.798460
K neighbours	0.695025	0.596962
Support vector machine	0.670948	0.615647
Decision tree	0.781711	0.790845
Random forest classifier	0.800980	0.813358

```

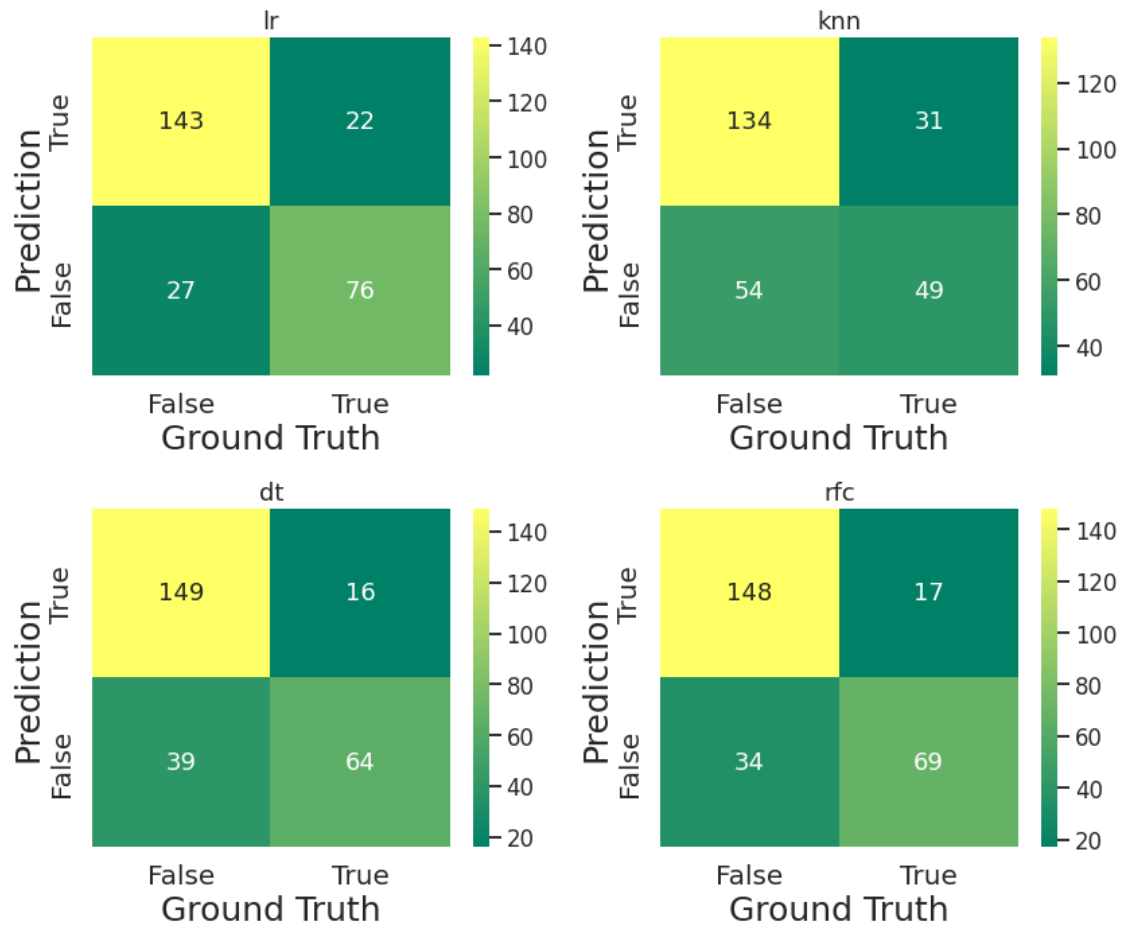
[198]: fig, axList = plt.subplots(nrows=2, ncols=2)
axList = axList.flatten()
fig.set_size_inches(12, 10)

models = coeff_labels = ['lr', 'knn', 'dt', 'rfc']
cm = [confusion_matrix(y_test, pred_lr),
      confusion_matrix(y_test, pred_knn),
      confusion_matrix(y_test, pred_dt),
      confusion_matrix(y_test, pred_rfc)]
labels = ['False', 'True']

for ax, model, idx in zip(axList, models, range(0,4)):
    sns.heatmap(cm[idx], ax=ax, annot=True, fmt='d', cmap='summer');
    ax.set(title=model);
    ax.set_xticklabels(labels, fontsize=20);
    ax.set_yticklabels(labels[::-1], fontsize=20);
    ax.set_ylabel('Prediction', fontsize=25);
    ax.set_xlabel('Ground Truth', fontsize=25)

plt.tight_layout()

```



Thus gonna select Logistic regression as my final model

11 Conclusion

In these projects we were able to create 5 different models to classify if a person survived or not
The final model had the following stats

[200]: lr_stats

```
[200]: precision    0.815939
       recall      0.817164
       accuracy    0.817164
       f1score     0.816255
       auc         0.802265
       Name: Logistic Regression, dtype: float64
```

12 Next steps

We can look into feature importance and discard some of the feature to improve the prediction We can also apply boosting , bagging and Stackin to improve the performance