

TimeSeries

May 23, 2022

1 Time Series Forecasting

Dataset Introduction: This dataset used here is Time Series Forecasting with Yahoo Stock Price from kaggle. This dataset is historical stocks of yahoo finance corp. The dataset consists of following columns

1. Date - Trading Date
2. High - the high refers to the maximum prices in a given time period.
3. Low - the low refers to the minimum prices in a given time period.
4. Open - prices at which a stock began trading in the same period.
5. close - the prices at which a stock ended trading in the same period.
6. Volume - Volume is the total amount of trading activity
7. Adj close - Adjusted values factor in corporate actions such as dividends, stock splits, and new share issuance.

Problem statement - Forecast the close series using deep learning

1.1 Exploratory data analysis

```
[60]: import sys
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.simplefilter(action='ignore')
import pandas as pd
from datetime import datetime
import tensorflow as tf
import keras
from keras.models import Sequential
from keras.layers import Dense, SimpleRNN, LSTM, Activation, Dropout
import math
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error
from tensorflow.keras.optimizers import Adam
```

```
[2]: data = pd.read_csv('/content/drive/MyDrive/ibm/Projects/timeseriesanalysis/
↳ yahoo_stock(Time series).csv', sep=",")
```

```
data.head()
```

```
[2]:
```

	Date	High	Low	Open	Close	\
0	2015-11-23	2095.610107	2081.389893	2089.409912	2086.590088	
1	2015-11-24	2094.120117	2070.290039	2084.419922	2089.139893	
2	2015-11-25	2093.000000	2086.300049	2089.300049	2088.870117	
3	2015-11-26	2093.000000	2086.300049	2089.300049	2088.870117	
4	2015-11-27	2093.290039	2084.129883	2088.820068	2090.110107	

	Volume	Adj Close
0	3.587980e+09	2086.590088
1	3.884930e+09	2089.139893
2	2.852940e+09	2088.870117
3	2.852940e+09	2088.870117
4	1.466840e+09	2090.110107

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1825 entries, 0 to 1824
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        1825 non-null   object
1   High        1825 non-null   float64
2   Low         1825 non-null   float64
3   Open        1825 non-null   float64
4   Close       1825 non-null   float64
5   Volume      1825 non-null   float64
6   Adj Close   1825 non-null   float64
dtypes: float64(6), object(1)
memory usage: 99.9+ KB
```

```
[5]: data.dtypes
```

```
[5]: Date        object
     High        float64
     Low         float64
     Open        float64
     Close       float64
     Volume      float64
     Adj Close   float64
     dtype: object
```

```
[6]: data.isnull().sum()
```

```
[6]: Date      0
      High      0
      Low       0
      Open      0
      Close     0
      Volume    0
      Adj Close  0
      dtype: int64
```

```
[7]: data.Date=pd.to_datetime(data['Date'])
```

```
[8]: data.set_index('Date',inplace=True)
```

```
[9]: data.head()
```

```
[9]:
```

	High	Low	Open	Close	Volume \
Date					
2015-11-23	2095.610107	2081.389893	2089.409912	2086.590088	3.587980e+09
2015-11-24	2094.120117	2070.290039	2084.419922	2089.139893	3.884930e+09
2015-11-25	2093.000000	2086.300049	2089.300049	2088.870117	2.852940e+09
2015-11-26	2093.000000	2086.300049	2089.300049	2088.870117	2.852940e+09
2015-11-27	2093.290039	2084.129883	2088.820068	2090.110107	1.466840e+09


```
Adj Close
```

Date	
2015-11-23	2086.590088
2015-11-24	2089.139893
2015-11-25	2088.870117
2015-11-26	2088.870117
2015-11-27	2090.110107

```
[10]: data.describe()
```

```
[10]:
```

	High	Low	Open	Close	Volume \
count	1825.000000	1825.000000	1825.000000	1825.000000	1.825000e+03
mean	2660.718673	2632.817580	2647.704751	2647.856284	3.869627e+09
std	409.680853	404.310068	407.169994	407.301177	1.087593e+09
min	1847.000000	1810.099976	1833.400024	1829.079956	1.296540e+09
25%	2348.350098	2322.250000	2341.979980	2328.949951	3.257950e+09
50%	2696.250000	2667.840088	2685.489990	2683.340088	3.609740e+09
75%	2930.790039	2900.709961	2913.860107	2917.520020	4.142850e+09
max	3645.989990	3600.159912	3612.090088	3626.909912	9.044690e+09



```
Adj Close
```

count	1825.000000
mean	2647.856284
std	407.301177

```
min    1829.079956
25%    2328.949951
50%    2683.340088
75%    2917.520020
max    3626.909912
```

```
[13]: plt.plot(data.Close)
```


```
[13]: [<matplotlib.lines.Line2D at 0x7fb3fda02510>]
```



1.2 Stationary Check

```
[14]: data.Close.hist(bins=50)
```

```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb3fd956590>
```




It seems that Close series is non stationary. Lets check this further

Lets decompose this series using additive and multiplicative decomposition

1.2.1 Decomposition


```
[25]: from statsmodels.tsa.seasonal import seasonal_decompose
decompose_add=seasonal_decompose(data['Close'], model='additive')
plt.figure(figsize=(15,15))
plt.subplot(411)
plt.plot(data['Close'], label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(decompose_add.trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(decompose_add.seasonal, label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(decompose_add.resid, label='Residual')
plt.legend(loc='best')
```

```
[25]: <matplotlib.legend.Legend at 0x7fb3f8badb50>
```



```
[24]: decompose_add=seasonal_decompose(data['Close'], model='multiplicative')
plt.figure(figsize=(15,15))
plt.subplot(411)
plt.plot(data['Close'], label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(decompose_add.trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(decompose_add.seasonal, label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(decompose_add.resid, label='Residual')
plt.legend(loc='best')
```

[24]: <matplotlib.legend.Legend at 0x7fb3f8cc82d0>




We can conclude that additive decomposition explains the series better than multiplicative decompositions

1.2.2 Plots


AutoCorrelation

```
[21]: from statsmodels.graphics.tsaplots import plot_acf
      from statsmodels.graphics.tsaplots import plot_pacf
      plt.rc("figure", figsize=(10,5))
      plot_acf(data.Close)
      print()
```



Partial Autocorrelation

```
[22]: plt.rc("figure", figsize=(10,5))
      plot_pacf(data.Close)
      print()
```



Hypothesis Testing 1. H_0 = non-stationary type. 2. H_1 = stationary series

```
[23]: from statsmodels.tsa.stattools import adfuller
```

```
[26]: result = adfuller(data['Close'])  
print('p-value:' +str(result[1]))
```

p-value:0.7975646340657458

Thus we can conclude Close series is a non stationary series

1.3 Train and test split

```
[65]: data2 = data[['Close']]  
timesteps = 50  
train = data2[:len(data)-timesteps]['Close'].values  
test = data2[len(train):]['Close'].values  
train=train.reshape(train.shape[0],1)  
test=test.reshape(test.shape[0],1)  
sc = MinMaxScaler(feature_range= (0,1))  
train = sc.fit_transform(train)  
X_train = []  
y_train = []  
  
for i in range(timesteps, train.shape[0]):  
    X_train.append(train[i-timesteps:i,0])  
    y_train.append(train[i,0])  
  
X_train = np.array(X_train)  
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)  
y_train = np.array(y_train)  
print('Training input shape: {}'.format(X_train.shape))  
print('Training output shape: {}'.format(y_train.shape))
```

Training input shape: (1725, 50, 1)

Training output shape: (1725,)

```
[66]: inputs = data2[len(data) - len(test) - timesteps:]  
  
inputs = sc.transform(inputs)  
  
X_test = []  
  
for i in range(timesteps, 100):  
    X_test.append(inputs[i-timesteps:i,0])  
  
X_test = np.array(X_test)
```

```
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], 1)
X_test.shape
```

[66]: (50, 50, 1)

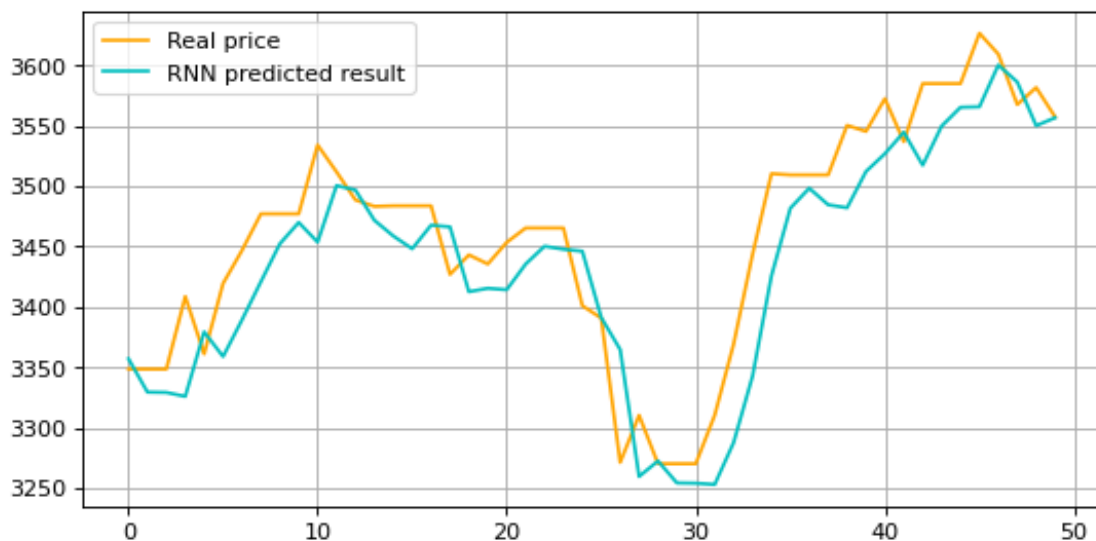
2 Models

2.1 RNN

```
[ ]: model=Sequential()
model.add(SimpleRNN(50,return_sequences=True, activation='relu',
    ↪input_shape=(X_train.shape[1],1)))
model.add(SimpleRNN(50,return_sequences=False,activation='relu'))
model.add(Dense(100))
model.add(Dense(25))
model.add(Dense(1))
opt1=Adam(learning_rate=1e-4,beta_1=0.9,beta_2=0.7)
model.compile(loss='mean_squared_error', optimizer=opt1)
model.fit(X_train, y_train, epochs=100, batch_size=32)
```

```
[68]: predicted = model.predict(X_test)
predicted = sc.inverse_transform(predicted)

plt.figure(figsize=(8,4), dpi=80, facecolor='w', edgecolor='k')
plt.plot(test,color="orange",label="Real price")
plt.plot(predicted,color="c",label="RNN predicted result")
plt.legend()
plt.grid(True)
plt.show()
```




3 LSTM

```
[ ]: model2 = Sequential()
model2.add(LSTM(70, input_shape=(X_train.shape[1],1)))
model2.add(Dense(1))

model2.compile(loss='mean_squared_error', optimizer='adam')
model2.fit(X_train, y_train, epochs=100, batch_size=32)
```

```
[74]: predicted2 = model2.predict(X_test)
predicted2 = sc.inverse_transform(predicted2)

plt.figure(figsize=(8,4), dpi=80, facecolor='w', edgecolor='k')
plt.plot(test,color="orange",label="Real price")
plt.plot(predicted2,color="c",label="LSTM model 1 Predictions")
plt.legend()
plt.grid(True)
plt.show()
```



```
[ ]: model3=Sequential()
model3.add(LSTM(50,return_sequences=True, activation='relu',_
↪input_shape=(X_train.shape[1],1)))
model3.add(LSTM(50,return_sequences=False,activation='relu'))
model3.add(Dense(100))
```

```

model3.add(Dense(25))
model3.add(Dense(1))
opt1=Adam(learning_rate=0.001,beta_1=0.9,beta_2=0.999)
model3.compile(loss='mean_squared_error', optimizer=opt1)
model3.fit(X_train, y_train, epochs=100, batch_size=10)


```

```

[75]: predicted3 = model3.predict(X_test)
predicted3 = sc.inverse_transform(predicted3)

plt.figure(figsize=(8,4), dpi=80, facecolor='w', edgecolor='k')
plt.plot(test,color="orange",label="Real price")
plt.plot(predicted3,color="c",label="LSTM model 2 Predictions")
plt.legend()
plt.grid(True)
plt.show()

```



```

[80]: print("rnn :"+str((mean_squared_error(y_train, model.predict(X_train)))))
print("lstm model 1 :"+str((mean_squared_error(y_train, model2.
    ↳predict(X_train)))))
print("lstm model 2 :"+str((mean_squared_error(y_train, model3.
    ↳predict(X_train)))))

```

```

rnn :0.00022443603655724226
lstm model 1 :0.00023869740249628116
lstm model 2 :0.0002348539642872087

```

4 Results

Among the 3 models all models performed moderately with rnn showing the least mse even though the lstm were trained longer

5 Next Steps

We can use hyperparameters tuning to improve accuracy of lstm and we need more data for lstm ,
We can also use hyperparameter tuning on rnn